

# COMPSCI 732

# SOFTENG 750

MongoDB

- MongoDB
  - Relational vs Document (“No-SQL”) databases
  - Examples
  - MongoDB Compass (DB browser tool)
- Mongoose: MongoDB from JavaScript applications
  - Modelling schema
  - Adding / removing / modifying documents



CS 732 / SE 750 – Module Four

# MongoDB

# Document databases

- Colloquially known as “No-SQL” databases
  - Because you don’t write SQL to interact with them!
- Originally designed to simplify the mapping between object-oriented business logic and the underlying data storage, without requiring ORM (Object Relational Mapping) middleware
- Individual documents (objects) are stored in various collections (similar to RDBMS tables) in a database
- Nested documents are trivial and don’t require separate collections (though sometimes you might want one)
- Relationships between different collections are easy
- Document / object properties can be indexed for increased speed / enforcing uniqueness

- Stores documents in a JSON-like structure
- Can interact with the DB via the command line using what is essentially JavaScript
- MongoDB Compass – a DB browser tool
- Can be local or cloud-based
- Wide community support

- Each MongoDB database can have an arbitrary number of **collections**
- Each *collection* can have any number of **documents**
- Each document can be represented as **JSON**
- Each document has a field called `_id`, which is an identifier that's guaranteed to be unique amongst all documents in its collection
- Documents may have any number of other fields

- There is **no requirement by default** for documents in a collection to conform to the same *schema* (i.e. have the same fields, etc).
- For example, it is valid for a single collection to contain all of these at the same time:

```
{
  _id: ObjectId("5ea10e4846c12e2d40b69a9a")
  name: "Thomas"
  age: 42
}
{
  _id: ObjectId("5ea11151d7deff616c570d5a")
  line1: "123 Some St"
  city: "Townsville"
}
{
  _id: ObjectId("5ea11562ccc5ba5378c9d96d")
  date: "2020-04-23T04:11:45.933+00:00"
  content: "This is a note."
}
```

In practice, developers will nearly always store objects of the same type in a single collection

- If we would like to enforce particular rules on inserted documents (e.g. required fields and data types of those fields, relationships between documents, etc...), we can use **schema validation**
- This can be done from the command line (<https://docs.mongodb.com/manual/core/schema-validation/>) or from Compass (see later slide)



- The following are valid datatypes for a field in a MongoDB document:
  - Commonly used: Binary, Boolean, Date, Decimal128, Double, Int32, Int64, ObjectId, String, Timestamp
  - Less common: Code, MaxKey, MinKey, BSONRegex, Symbol
- In addition, fields can be null, undefined, or other objects (essentially sub-documents)

# MongoDB – downloading

- Download from: <https://www.mongodb.com/>
- Can create an account using Google credentials
- Free for personal use
- Can download local version (Community Server, from the Software tab) and / or create a free Atlas Cloud database
- Can also download Compass, the DB browser tool

- MongoDB installations come with a *shell*
  - The mongo shell is a JavaScript REPL interface to a database
  - <https://docs.mongodb.com/manual/mongo/>
- We can also use the GUI – MongoDB Compass
- Finally, we can interact with the database through our own node.js code using mongoose.

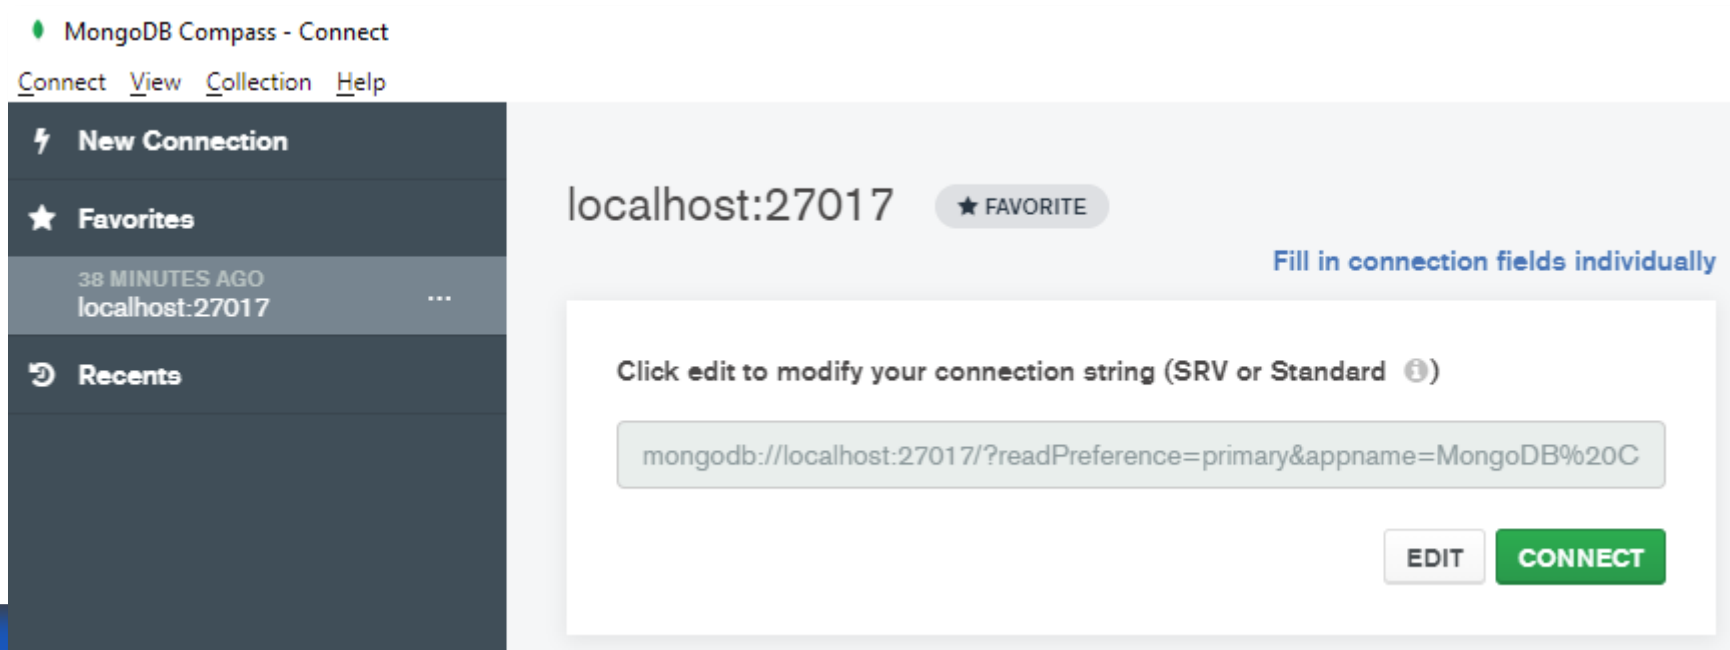
CS 732 / SE 750 – Module Four

# MongoDB Compass



# MongoDB Compass

- Once the tool is open, connect using your *connection string*
  - This is mongodb://localhost:27017/ by default for local installations
  - You can see your cloud connection string on the Atlas website



# MongoDB Compass

- Once connected we can see a list of databases on the connected server

HOST  
localhost:27017

CLUSTER  
Standalone

EDITION  
MongoDB 4.0.9 Community

Filter your data






> admin

> config

> local

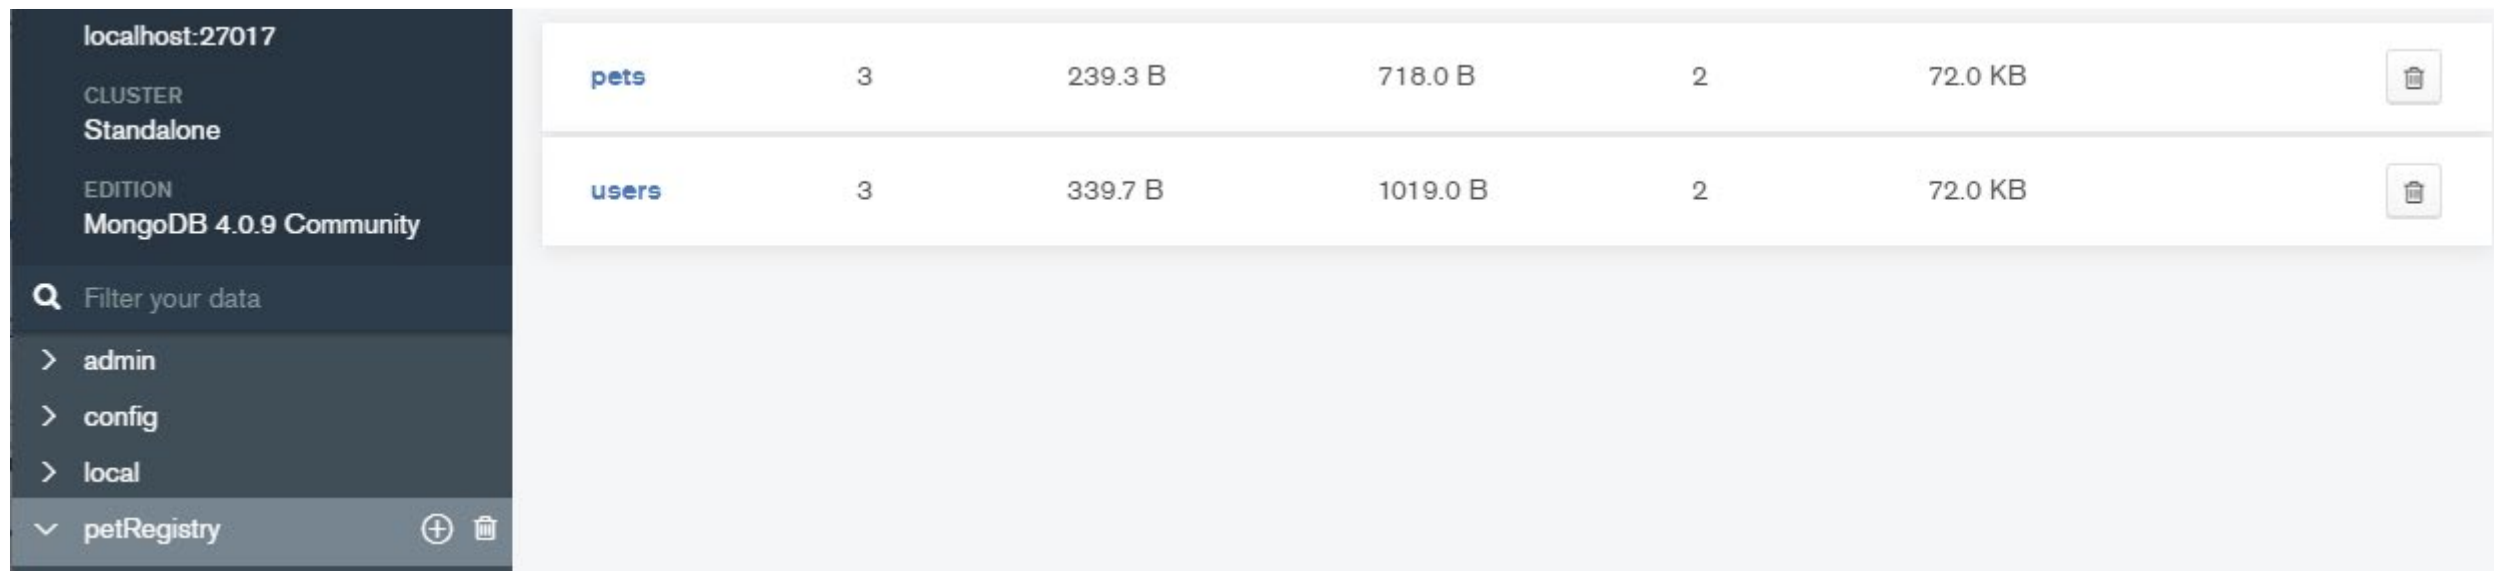
> petRegistry

> vly2



admin	16.0KB	0	1	
config	36.0KB	0	2	
local	44.0KB	1	1	
petRegistry	72.0KB	2	4	
vly2	48.0KB	2	4	

# MongoDB Compass

- Clicking on one will show us the collections in that database

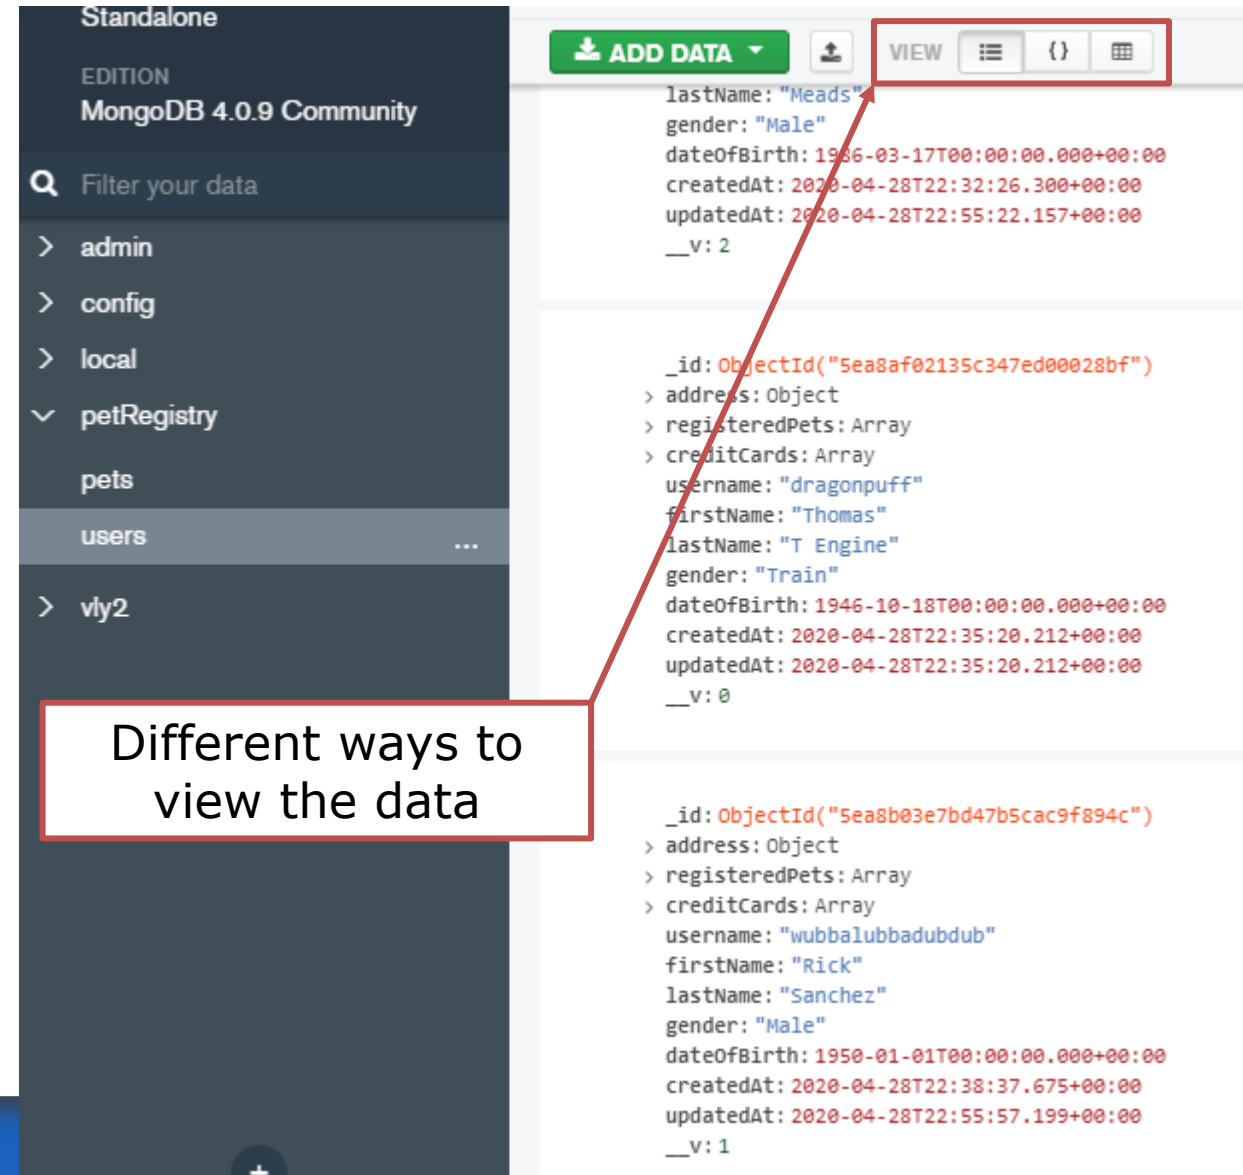


The screenshot shows the MongoDB Compass interface. On the left sidebar, the connection is 'localhost:27017', the cluster is 'Standalone', and the edition is 'MongoDB 4.0.9 Community'. Below this is a search bar 'Filter your data' and a list of databases: 'admin', 'config', 'local', and 'petRegistry' (which is expanded). The main panel displays the collections for the 'petRegistry' database. There are two collections: 'pets' and 'users'. Each collection row shows the number of documents, the size of the collection, the size of the index, the number of indexes, and the size of the index. Each row also has a trash icon for deleting the collection.

Collection	Documents	Collection Size	Index Size	Index Count	Index Size	Action
<b>pets</b>	3	239.3 B	718.0 B	2	72.0 KB	
<b>users</b>	3	339.7 B	1019.0 B	2	72.0 KB	


# MongoDB Compass

- Clicking on a collection will show us the documents in that collection
- From here we can add arbitrary documents to the collection using "Add Data"





- We can filter on this page too – i.e. query for documents which match certain criteria



- We can filter by match, exclusion, comparison, date, array contents, ...
  - <https://docs.mongodb.com/compass/current/query/filter/>

# MongoDB Compass

- We can edit individual documents
  - When hovering over a document, a context menu will appear showing the edit and delete buttons



- The database itself and its command-line and GUI tools have extensive documentation:
  - [MongoDB](#)
  - [Mongo shell](#)
  - [Compass](#)
- However, we will focus on interaction from our own JavaScript code...



CS 732 / SE 750 – Module Four

# Mongoose

- Mongoose allows us to interact with MongoDB databases via node.js (including the Express apps that often form the backends of our React webapps)
- Install as follows:

```
npm install mongoose
```

- Import:

```
import mongoose from 'mongoose';
```

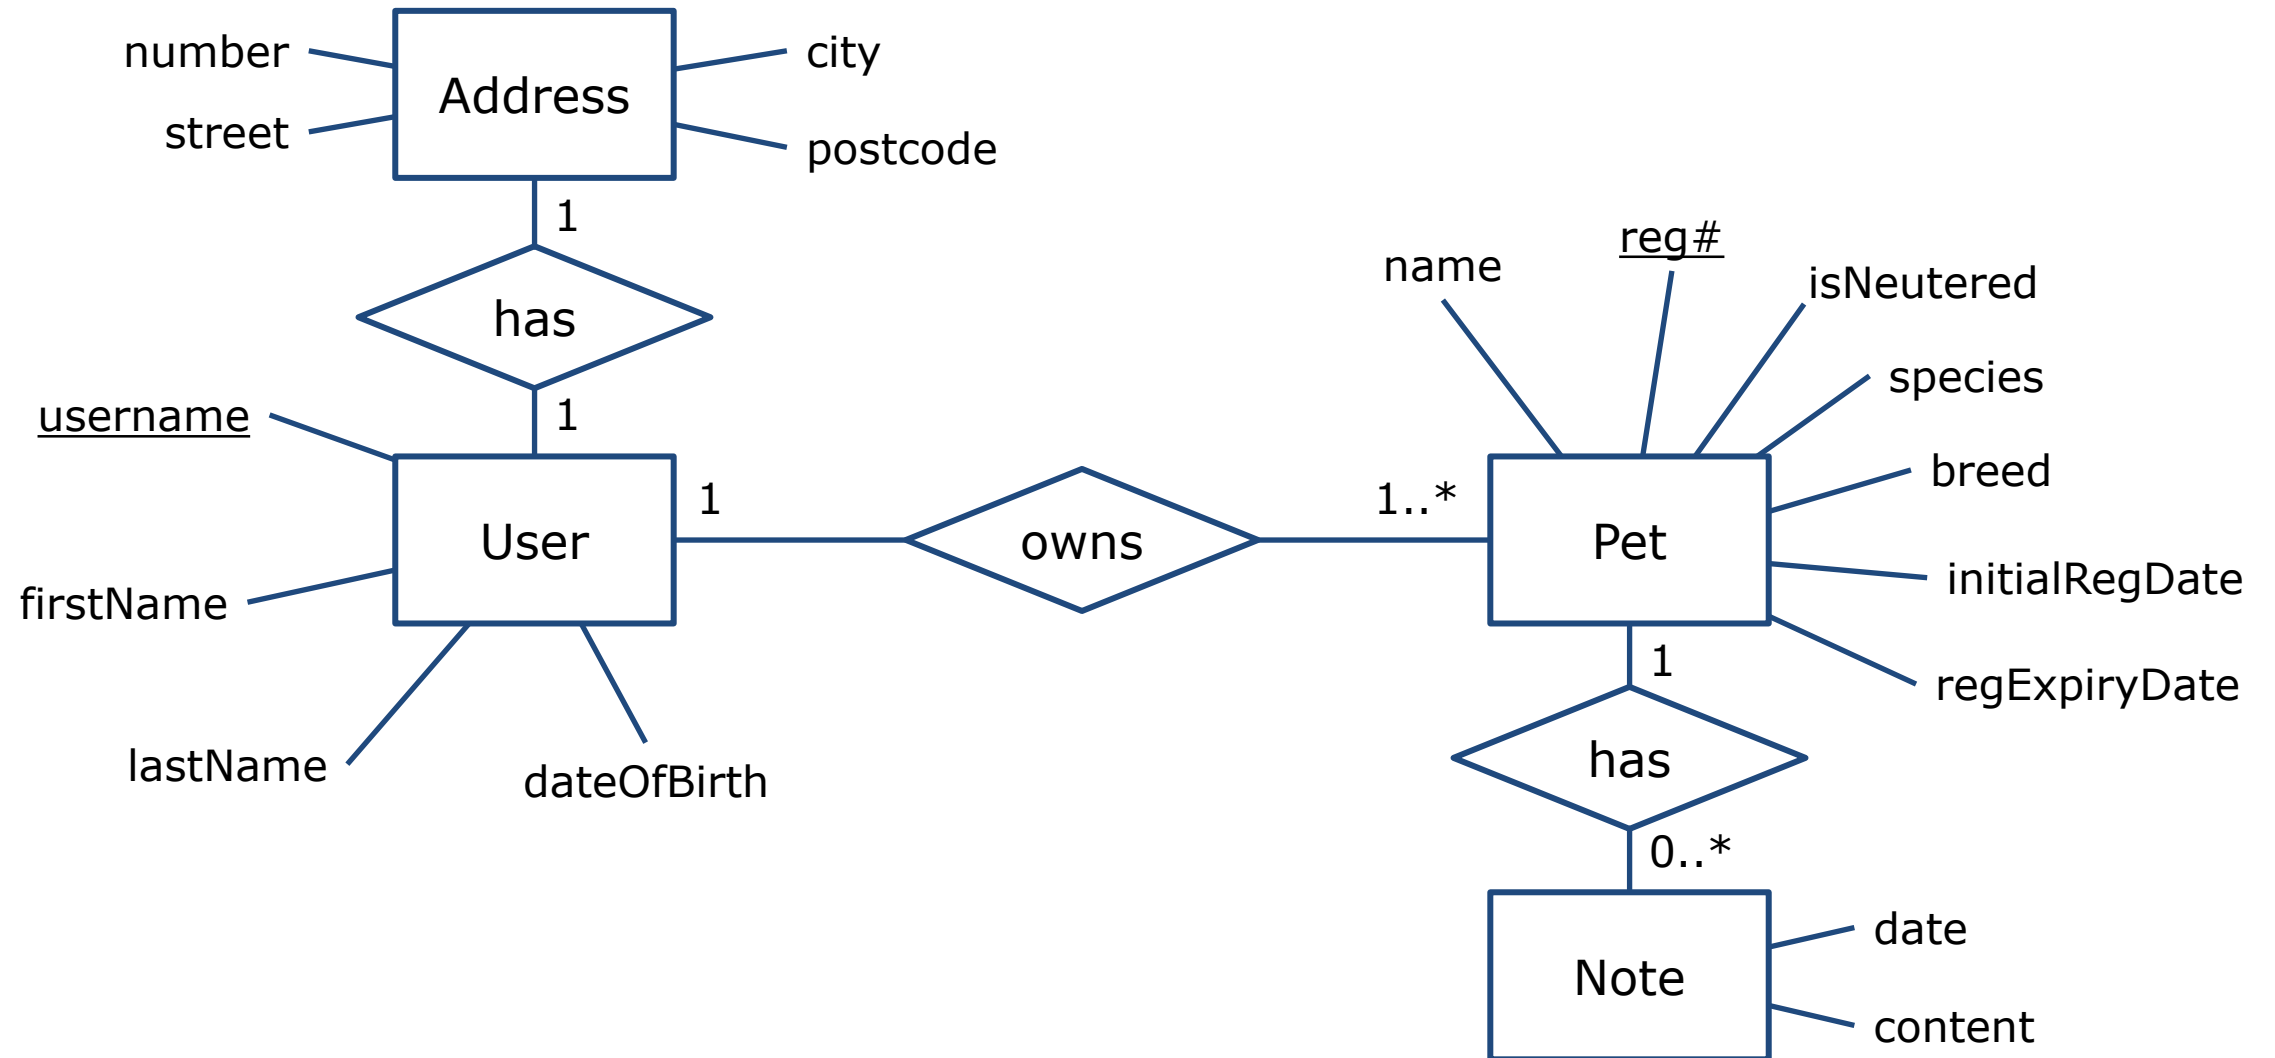


CS 732 / SE 750 – Module Four

# Mongoose schema definition

- Let's build a simple database for a pet registry webapp:
  - Users have a username, first & last names, date of birth, gender, and address
    - A user's address has several fields such as street name, suburb, etc.
  - A user may have any number of pets registered
  - Each pet has a registration number, a name, a breed & species, an initial registration date, a registration expiry date, and info on whether the pet has been neutered.
    - Any number of additional notes may be made about any pet. Each note has a date and content (text).

# Schema definition





- If we were building a traditional relational database (e.g. SQLite, MySQL, Postgres...), we might:
  - Form a relational model from our ER diagram
  - Use this to inform our SQL CREATE TABLE statements

# Possible relational model

User(username, firstName, lastName, dateOfBirth)

Address(number, street, city, postcode, ownerUsername)

Pet(regnum, name, isNeutered, species, breed, initialRegDate, expiryDate, ownerUsername)

Note(petRegnum, date, content)

# Possible relational model

User(id, username, firstName, lastName, dateOfBirth)

Address(number, street, city, postcode, owner\_id)

Pet(id, regnum, name, isNeutered, species, breed, initialRegDate, expiryDate, owner\_id)

Note(pet\_id, date, content)

# Possible relational model

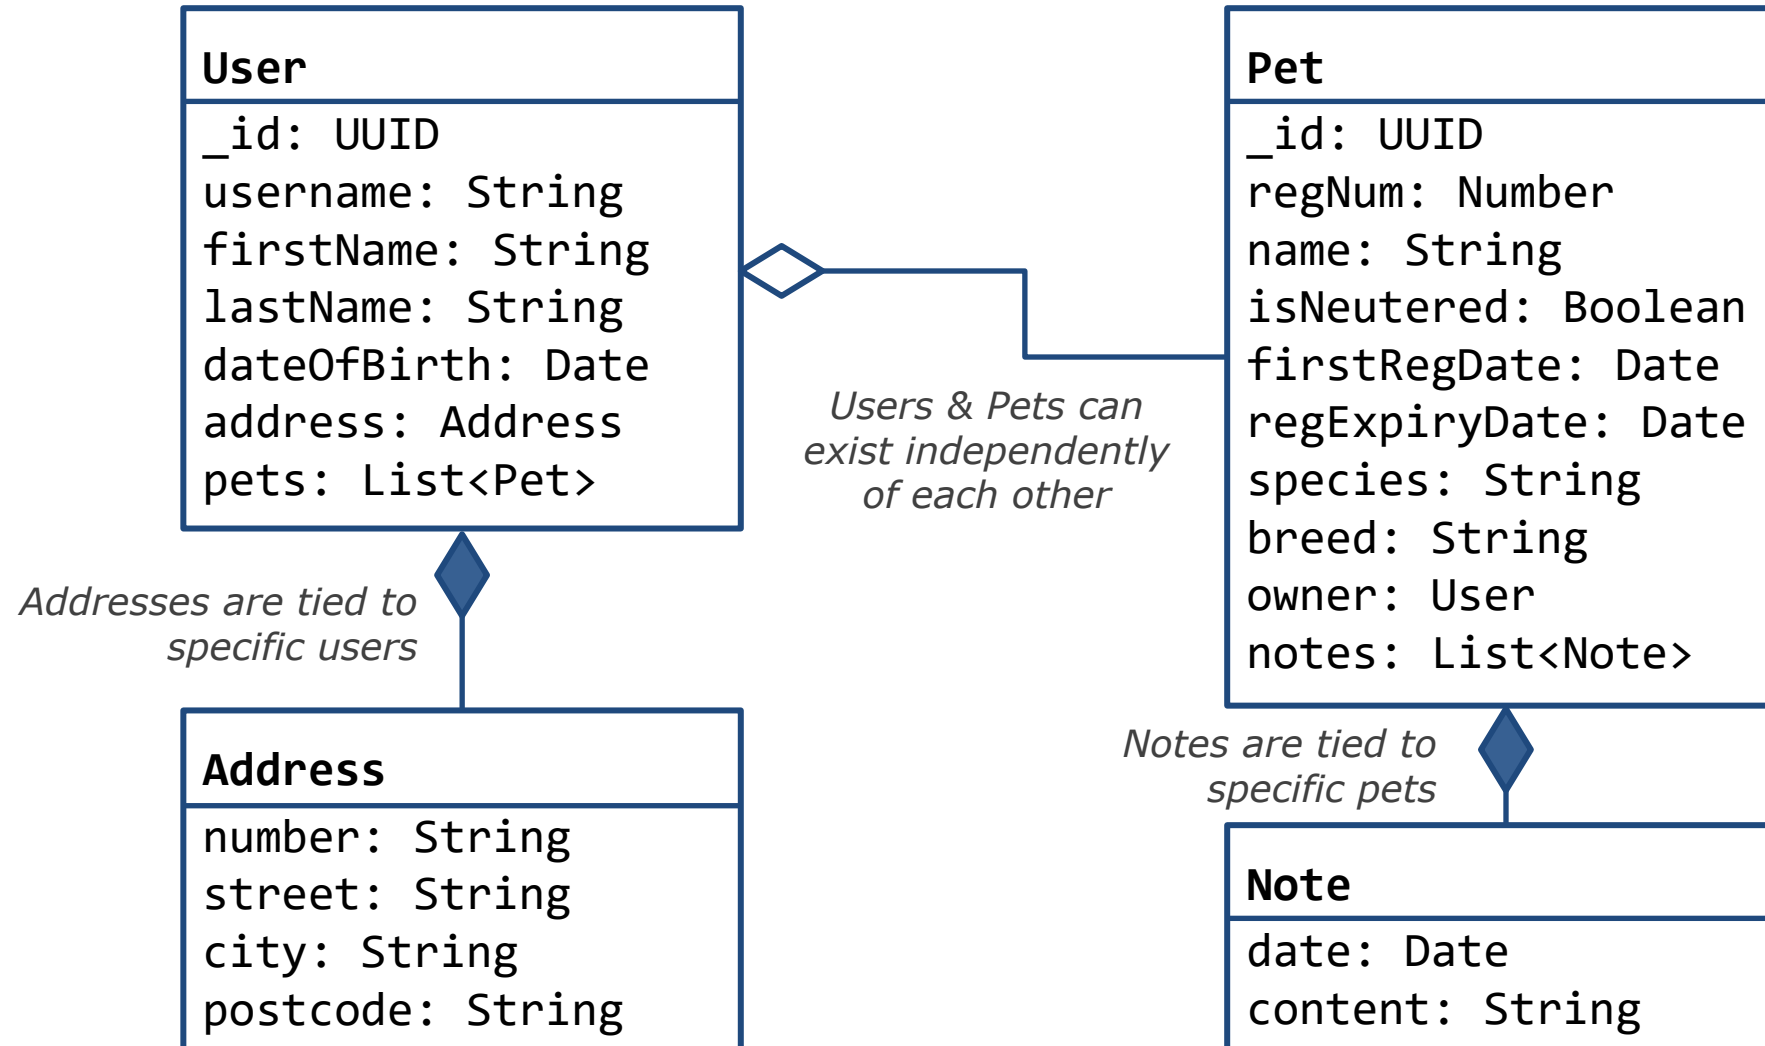
User(id, username, firstName, lastName, dateOfBirth,  
address\_number, address\_street, address\_city, address\_postcode)

Pet(id, regnum, name, isNeutered, species, breed,  
initialRegDate, expiryDate, *owner\_id*)

Note(*pet\_id*, date, content)

- When designing a MongoDB database schema, it might help more to think of our system in terms of **objects** and **classes** – similar to object-oriented software design
- The classes might inform the design of our schema, while individual objects may be documents in our database

# Object model



# Potential queries

- When designing a schema, it can be useful to think of the queries one might ask of the database. In this case:
  - Get all Users
  - Get all Pets
  - Get a User with a particular username
  - Get a Pet with a particular registration number
  - Get all Pets owned by a particular User
  - Get the User who owns a particular Pet
- Our **Users** and **Pets** are important entities which need to be queried independently of each other – they can be stored in their own **collections**
- **Addresses** and **Notes** can't exist independently. They can be **sub-documents**

# Schema definition in mongoose

- To define the schema for a particular type of document, we create an instance of a Schema object
- We can then use the `mongoose.model()` function to create a class definition for our document type, which we can then create instances of.

```
import mongoose from 'mongoose';  
const Schema = mongoose.Schema;
```

```
const userSchema = new Schema({ ... });
```

```
const User = mongoose.model('User', userSchema);
```

```
const user = new User();
```

Create a new Schema which defines what a User looks like in our system

Create the User class. The name given as a string will also determine the name of the MongoDB collection used (it will use the plural form of the given word, in all lower case – e.g. users in this case).

Create an individual User instance



# Schema definition in mongoose

```
const userSchema = new Schema({
```

```
  username: { type: String, unique: true },  
  firstName: String,  
  lastName: String,  
  gender: String,  
  dateOfBirth: Date,
```

```
  address: {  
    line1: String,  
    line2: String,  
    suburb: String,  
    city: String,  
    postcode: Number
```

```
  },
```

```
  creditCards: [{ lastFourDigits: String, encryptedInfo: String }],
```

```
  registeredPets: [{ type: Schema.Types.ObjectId, ref: 'Pet' }],
```

```
}, {  
  timestamps: {}  
});
```

Define a field which must be unique amongst all Users. This will be indexed and uniqueness will be enforced in MongoDB.

Standard fields

Address is an “embedded document”. It won’t be a separate collection. Each User will have an address field which in turn will have these fields.

The array [] syntax means that Users will have *an array of* credit cards, each one with the two fields given here

Each element in the registeredPets array should be an ObjectId referring to the `_id` of a document in the pets collection

# Schema definition in mongoose

```
const userSchema = new Schema({  
  
  username: { type: String, unique: true },  
  firstName: String,  
  lastName: String,  
  gender: String,  
  dateOfBirth: Date,  
  
  address: {  
    line1: String,  
    line2: String,  
    suburb: String,  
    city: String,  
    postcode: Number  
  },  
  
  creditCards: [{ lastFourDigits: String, encryptedInfo: String }],  
  registeredPets: [{ type: Schema.Types.ObjectId, ref: 'Pet' }],  
}, {  
  timestamps: {}  
});
```

The second object supplied here contains additional config info.

Here, we can specify additional options such as auto-generating createdAt and updatedAt fields (as shown here), removing the auto-generated \_id, or many other options.

See: <https://mongoosejs.com/docs/guide.html>

# Schema definition in mongoose

- In addition to fields and sub-documents, we can also add extra:
  - Instance methods

```
const animalSchema = new Schema({ name: String, type: String });
```

```
animalSchema.methods.findSimilarTypes = function (cb) {  
  return this.model('Animal').find({ type: this.type }, cb);  
};
```

```
const Animal = mongoose.model('Animal', animalSchema);  
const dog = new Animal({ type: 'dog' });
```

```
dog.findSimilarTypes(function (err, dogs) {  
  console.log(dogs); // woof  
});
```

# Schema definition in mongoose

- In addition to fields and sub-documents, we can also add extra:
  - Static methods

```
const animalSchema = new Schema({ name: String, type: String });
```

```
animalSchema.statics.findByName = function (name) {  
  return this.find({ name: new RegExp(name, 'i') });  
};
```

```
const Animal = mongoose.model('Animal', animalSchema);
```

```
let animals = await Animal.findByName('fido');
```

# Schema definition in mongoose

- In addition to fields and sub-documents, we can also add extra:

- Virtuals ([JavaScript link](#))

```
const userSchema = new Schema({firstName: String, lastName: String});
```

```
userSchema.virtual('fullName')  
  .get(function () { return `${this.firstName} ${this.lastName}`; })  
  .set(function (value) {  
    this.firstName = value.substr(0, value.indexOf(' '));  
    this.lastName = value.substr(value.indexOf(' ') + 1);  
  });
```

```
const User = mongoose.model('User', userSchema);
```

```
const axl = new User({ firstName: 'Axl', lastName: 'Rose' });
```

```
console.log(axl.fullName);
```

```
axl.fullName = 'William Rose';
```

```
console.log(axl.firstName);
```

← "Axl Rose"

← "William"



CS 732 / SE 750 – Module Four

# Mongoose schema validation

# Schema validation

- We can further extend our schema with validation logic which will be run when we try to save an object to the database
  - Objects failing the validation check will not be saved
- We can also run the validation logic ourselves in an attempt to pre-empt / fix issues before trying to save

```
const meow = new Schema({  
  name: { type: String, required: true }  
});  
const Cat = db.model('Cat', meow);
```

The name field is required.

```
const cat = new Cat();  
cat.save(function (error) {  
  assert.equal(error.errors['name'].message, 'Path `name` is required.');
```

Trying to save a cat with no name will give us an error

```
});  
  
const error = cat.validateSync();  
assert.equal(error.errors['name'].message, 'Path `name` is required.');
```

We can use `validateSync()` to run the validation logic whenever we like

# Schema validation

```
const breakfastSchema = new Schema({
  eggs: {
    type: Number,
    min: [6, 'Too few eggs'],
    max: 12
  },
  bacon: {
    type: Number,
    required: [true, 'Why no bacon?']
  },
  drink: {
    type: String,
    enum: ['Coffee', 'Tea'],
    required: function () {
      return this.bacon > 3;
    }
  }
});
```

If we have eggs (it's not a required field), we must have between 6 – 12. "Too few eggs" is a custom error message that we'll get if the min validation fails.

We must have some bacon for breakfast. "Why no bacon?" is a custom error message we'll get if there's no bacon.

If we have a drink, it must be either Coffee or Tea.

We must have a drink if we have more than 3 strips of bacon.



- We can define custom validator functions for arbitrary validation logic

```
const userSchema = new Schema({  
  phone: {  
    type: String,  
    validate: {  
      validator: function (v) {  
        return /\d{3}-\d{3}-\d{4}/.test(v);  
      },  
      message: props => `${props.value} is not a valid phone number!`  
    }  
  }  
});
```

A user's phone number must match this regular expression (XXX-XXX-XXXX, where X's are digits)

This is the custom error message we'll get if this validation fails



CS 732 / SE 750 – Module Four

# Working with documents

# Modelling and object (document) creation

- Once we've modelled our schema:

```
const User = mongoose.model('User', userSchema);
```

- We can then create instances of the model class, which will correspond to MongoDB documents:

```
const anne = new User();  
anne.fullName = 'Anne Hathaway';  
anne.dateOfBirth = new Date('1982-11-12');
```

← Can get / set fields on an individual basis...

```
const bob = new User({  
  firstName: 'Bob', lastName: 'Ross',  
  dateOfBirth: new Date('1942-10-29')  
});
```

← And / or supply values in the constructor as so

```
console.log(bob._id);
```

The `_id` field is auto-generated; we didn't have to define this. If we need to, we can disable this behaviour

(see: [https://mongoosejs.com/docs/guide.html#\\_id](https://mongoosejs.com/docs/guide.html#_id))

# Connecting, saving & deleting

- Many of these functions are *async* and thus return [promises](#). We can await them as shown here from within our own async functions, or use any other promise management syntax.

- To connect to a database:

```
mongoose.connect('mongodb://localhost:27017/petRegistry',  
  { useNewUrlParser: true });
```

- To save a document instance:

```
const bob = new User({ ... });  
await bob.save();
```

- To delete an instance:

```
await User.deleteOne({ username: 'anhydrous' });
```

Deletes the first document in the users collection matching the given criteria. See “querying” for further discussion of valid criteria. There is also a deleteMany() function.



CS 732 / SE 750 – Module Four

# Querying in mongoose

# Querying in mongoose

- We can conduct queries using the following static methods on our model classes (those generated with `mongoose.model`):
  - `findById()`: finds and returns the document with the given `_id`, if any
  - `findOne()`: finds and returns the first document matching the given criteria, if any
  - `find()`: finds and returns an array of all documents matching the given criteria
- All such methods are async and thus can be awaited – but **are not promises**.

*Find the user whose `_id` is `5ea8aede135c347ed00028be`*

```
const user = await User.findById('5ea8aede135c347ed00028be');
```

*Find all users*

```
const allUsers = await User.find();
```

# Querying in mongoose

*Find the first user whose username is anhydrous*

```
const user = await User.findOne({ username: 'anhydrous' });
```

*Find all users whose age is greater than 42*

```
const users = await User.find({ age: { $gt: 42 } });
```

*Find all users whose age is greater than 42, and less than or equal to 100*

```
const users = await User.find({ age: { $gt: 42, $lte: 100 } });
```

*Find all users whose city (part of their address) is either Auckland or Hamilton*

```
const users = await User.find({ 'address.city': { $in: ['Auckland', 'Hamilton'] } });
```

*Find all users whose address line 1 contains the text "Some Street", AND who have at least one registered pet*

```
const users = await User.find({  
  'address.line1': /Some Street/,  
  registeredPets: { $not: { $size: 0 } }  
});
```

For more query operators, see:  
<https://docs.mongodb.com/manual/reference/operator/query/>

# Populating fields referencing other collections

- When we have a field that's an ObjectId or an array of ObjectIds referencing another collection, we can populate that field with the data from the other collection using the populate method.
  - This works at query-time:

```
const user = await User
  .findOne({ username: 'anhydrous' })
  .populate('registeredPets');
```

This will make the registeredPets field of the returned user be an array of Pet objects, rather than an array of ObjectIds.
  - Or at any point afterwards.

```
const pets = await Pet.find({ ... });
await Pet.populate(pets, 'owner');
```

This will *convert* the owner field of all pets in the given array to the matching User object (or null if there's no matching user)





CS 732 / SE 750 – Module Four

# Building a REST API for a MongoDB app

# RESTful APIs for MongoDB applications

- Using mongoose, we can easily integrate MongoDB into our node.js apps
  - Including Express web apps / services
- Using the Express knowledge we already have from Week Three, plus the mongoose knowledge from these slides, we can build an API which utilizes MongoDB
- Let's use our running Articles app as an example...

# An Articles REST API backed by MongoDB

- Defining the schema

```
const Schema = mongoose.Schema;

const articleSchema = new Schema({
  title: { type: String, required: true },
  date: Date,
  image: String,
  content: String
}, {
  timestamps: {}
});

const Article = mongoose.model('Article', articleSchema);
```

# An Articles REST API backed by MongoDB

- Create

```
async function createArticle(article) {  
  
    const dbArticle = new Article(article);  
    await dbArticle.save();  
    return dbArticle;  
}  
  
router.post('/', async (req, res) => {  
    const newArticle = await createArticle(req.body);  
  
    if (newArticle) return res.status(HTTP_CREATED)  
        .header('Location', `/api/articles/${newArticle._id}`)  
        .json(newArticle);  
  
    return res.sendStatus(422);  
})
```

# An Articles REST API backed by MongoDB

- Retrieve One

```
async function retrieveArticle(id) {  
  return await Article.findById(id);  
}  
  
router.get('/:id', async (req, res) => {  
  const { id } = req.params;  
  
  const article = await retrieveArticle(id);  
  
  if (article) return res.json(article);  
  return res.sendStatus(HTTP_NOT_FOUND);  
});
```

# An Articles REST API backed by MongoDB

- Retrieve All

```
async function retrieveArticleList() {  
    return await Article.find();  
}  
  
router.get('/', async (req, res) => {  
    res.json(await retrieveArticleList());  
});
```

# An Articles REST API backed by MongoDB

- Update

```
async function updateArticle(article) {  
  
    const dbArticle =  
        await Article.findOneAndUpdate({ _id: article._id }, article);  
    return dbArticle !== undefined;  
}  
  
router.put('/:id', async (req, res) => {  
    const { id } = req.params;  
    const article = req.body;  
    article._id = id;  
    const success = await updateArticle(article);  
    res.sendStatus(success ? HTTP_NO_CONTENT : HTTP_NOT_FOUND);  
});
```

# An Articles REST API backed by MongoDB

- Delete

```
async function deleteArticle(id) {  
  await Article.deleteOne({ _id: id });  
}  
  
router.delete('/:id', async (req, res) => {  
  const { id } = req.params;  
  await deleteArticle(id);  
  res.sendStatus(HTTP_NO_CONTENT);  
});
```



# Full stack!

- Combining this REST API, backed by a database, with our React knowledge... we have now mastered the MERN stack!



- All frameworks & libraries covered in these slides are far more comprehensive than we can cover in one lecture:
- All have excellent documentation, which is a great starting point for further reading:
  - [MongoDB](#)
  - [Mongoose](#)